

Case Study: Regression

Machine Learning



Course Outline

Assignments are submitted at the same date as the next assignment comes out

Wk	Lecture 1	Lecture 2	Individual Assignment	Project Assignment
1	Case Study: Regression	Case Study: Classification	A1: Car Price Prediction	
2	Batch Gradient Descent	Stochastic / Mini-Batch Gradient Descent		
3	Regularization	Binary Logistic Regression	A2: Car Price Prediction II	
4	Multinomial Logistic Regression	Gaussian Naive Bayes		
5	Multinomial Naive Bayes	K-Nearest Neighbors / Support Vector Machine	A3: Car Price Prediction III	
6	Decision Tree	Bagging / Random Forests / Ada Boosting / Gradient Boosting		
7	No class	Midterm Exam		Phase 1: Reading paper round 1 (KDD)
8	K-Means Clustering	Gaussian Mixture		Phase 1: Reading paper round 2 (KDD)
9	Principal Component Analysis	PyTorch Linear Regression		Phase 2: Proposal - Paper writing (Intro, Related Work, Method)
10	PyTorch Logistic Regression	Project Proposal Presentation		
11	Convolutional Neural Network	Recurrent Neural Network		Phase 3: Experiment - Paper writing (Intro, Related Work, Method, Results)
12	Reinforcement Learning	Project Time		
13	Project Time	Project Progress Presentation		
14	Project Time	Project Time		Phase 4: Conclusion - Paper writing (Abstract, Intro, Related Work, Method, Results, Discussion, Conclusion)
15	No class	Final Exam		
16	No class	Final Project Presentation		



Background - Regression

- Predict **blood pressure level** using **BMI**?
- Predict **GDP** using **egg price, gold price, oil price**?
- Predict **life expectancy** using **population, GDP**?

We called population, GDP as **features** (a.k.a predictors, independent variables), and life expectancy as **labels** (a.k.a targets, dependent variables)

Notice all the **labels** here are **continuous** values?

Regression is a **supervised** algorithm to *predict* **continuous** values

- Labels must be continuous; features can be categorical or continuous
- Both features and labels can be univariate or multivariate
- *Supervised* - has both features and labels; *Unsupervised* - only has features
- *Continuous* - opposite of categorical variables, e.g., gold price is continuous while gender is categorical
 - Continuous value can be change to categorical variables by binning!



Big picture of ML

Remember: Splitting > Imputation > Scaling

Step 2 to 5 - can be done iteratively and order may swap, it's ok

1. Load data -> 2. Exploratory Data Analysis -> 3. Feature engineering -> 4. Feature selection -> 5. Preprocessing

- CSV, JSON, Database
- Renaming
- Label encoding
- Train / dev / test set

- Check class imbalance
- Countplot
- Distribution plot
- Boxplot
- Scatter plot
- Correlation Matrix
- Predictive Power Score

- Dimensionality reduction
- Feature splitting (e.g., date)
- Creating features (e.g., some equation)

- Select your X (features) and y (target)
- In ML, it's better to choose X
- In DL, we usually just input all features

- Null values
- Outliers
- Fix class imbalance (up-/down-sampling)
- Typos / Entry errors / Duplicates / IDs
- Scaling (min-max; standardize)

6. Model selection -> 7. Testing -> 8. Analysis -> 9. Inference -> 10. Deployment

Supervised

- Regression
- Classification

Unsupervised

- Clustering
- Dimensionality reduction

Reinforcement

- PPO
- Q-learning

Apply your best model on your test set

Analyze your model, e.g., **feature importance**

Apply your best model on some unseen data, and see whether it makes sense

• HTML/CSS/JS

- Flask
- Django
- FastAPI

• Docker

METRIC

- **Regression** (r^2 , MSE)
- **Classification** (recall, precision, f1)
- **Clustering** (inertia)
- **Dimensionality reduction** (mean squared distance between the original data and the reconstructed data)
- **Reinforcement learning** (cumulative rewards)

- Overfitting
- Underfitting

- MLFlow
- Wandb
- Tensorboard

- **Cross-validation**
 - With/Without Sampling
- **Grid search**

TOOLS

- **Python**, R - programming tool
 - **NumPy** (matrix manipulation), **Pandas** (Excel-like), **Matplotlib/Seaborn** (visualization), **Sklearn** (machine learning), **PyTorch** (deep learning), **Huggingface** (models)
- **Tableau**, **Power BI** - Business Intelligence (BI) tools
- **Microsoft Azure**, **Rapidminer**, **Weka** - data science and machine learning tools
- **SPSS**, **SAS**, **JASP** - statistical tool

TOP VENUES

1. ML (KDD)
2. DL (ICML, NIPS, ICMR)
3. NLP (ACL, EMNLP)
4. CV (CVPR, ICCV)



1. Load data

- Getting data
 - www.kaggle.com/datasets
 - API, e.g., [Instagram API](#), [Twitter API](#), [Weather API](#), [Stock Price API](#)
 - [Google Dataset search](#)
 - Collect data using sensors
- Data format: CSV, JSON, Databases (SQL)

```
df = pd.read_csv('data/Life_Expectancy_Data.csv')
```

Figure: Python can easily read your csv file



2. Exploratory Data Analysis

- Perhaps the most important step - understand your data
- **Label encoding:** encode so we can visualize properly, as well as performing correlation etc.
 - categorical variables into integers, e.g., male => 0; female => 1
- Univariate
 - Countplot
 - Distribution plot
- Multivariate
 - Boxplot
 - Scatter plot
 - Correlation Matrix
 - Predictive Power Score



One-hot encoding

- Recall label encoding which we turn categories into 0, 1, 2 etc.
- When we have **more than two categories**, if we encode into 0, 1, 2
 - we create a unintentional order, i.e., the model "may" think that $0 < 1 < 2$
- Possible solution: **one hot encoding**
 - E.g., Male, Female, Unknown \Rightarrow [1, 0, 0] if male; [0, 1, 0] if female
 - Limitation
 - what if we have like 5000 categories....
 - one hot encode this will result in 5000 columns --> too much! -> Two choices:
 - Group these categories into bigger categories, and then one-hot encode
 - Do label encoding anyway.....but note the possible order effect
- **Tips:** one thing you need to know is that you can always cut down one column
 - [1, 0, 0], [0, 1, 0], [0, 0, 1] is same as [1, 0], [0, 1], [0, 0] by setting `'drop_first=True'`



3. Feature engineering

- **Dimensionality reduction** - reduce dimensions, e.g, Principal Component Analysis (maximizes variances), Discriminant Analysis (maximizes separability)
- **Feature splitting** - Jan 26, 2023 ==> Monday or January
- **Creating features (e.g., some equation)** - combining the total sales from each sale department

Note: After such engineering, you can always go back to exploratory analysis



4. Feature selection

- Choose the most salient X
 - Rule of thumb: Good features MUST NOT BE correlated, i.e., independent
 - Rule of thumb: Correlation is not causation; don't pick features using correlation only; it should make sense!
 - E.g., *Number of trees is correlated with number of divorce*
 - Rule of thumb: For ML, less features are usually better (but NOT necessarily for DL)
- Specify the y
- Split train / test
 - Rule of thumb: Always split BEFORE preprocessing, to prevent data leakage
 - Can be done in this order: (1) splitting, (2) imputation, (3) scaling



5. Preprocessing

Numbers: 4, 5, 6, 8, 8, 9, 10, 11, 13, 18
 50% = $(8 + 9) / 2 = 8.5$
 75% = $(10, 11, 13, 18) / 2 = 12$
 25% = $(4, 5, 6, 8) / 2 = 5.5$
 IQR = $Q3 - Q1 = 12 - 5.5 = 6.5$
 Min = $Q1 - 1.5(IQR) = 5.5 - 1.5(6.5) = -4.25$
 Max = $Q3 + 1.5(IQR) = 12 + 1.5(6.5) = 21.25$


● Imputing missing values

- For missing “numbers”,
 - Replace with average | median | regression results | 0
- For missing “categories”
 - Replace with mode | 5 M, 6 F == 5/6 | “No category”
- 1st rule of thumb: replacing should not drastically change the distribution
- 2nd rule of thumb: don't delete if you are not sure what you are dealing with
- 3rd rule of thumb: replace ONLY features, NOT labels. Don't create fake labels.

● Outliers

- Use visualization such as box plot or z-score
- Rule of thumb: Don't just delete using box plot rule. Use your common sense to understand what is possible, errors, and impossible

● Typos / Entry errors / Duplicates / IDs



5. Preprocessing - scaling

- Scale your features
 - Help the model to learn faster
 - **Standardization**
 - $(x - \text{mean}) / \text{std}$
 - when your data follows normal distribution
 - **Normalization**
 - $(x - x_{\min}) / (x_{\max} - x_{\min})$
 - when your data DOES NOT follow normal distribution (e.g., audio, signal, image)
- Rule of thumb: Scale after you split and fill all the missing values
- Rule of thumb: Scale your test set using training distribution, NOT testing distribution
- Rule of thumb: DON'T scale your categorical features



6. Model selection - many algorithms

- Compare algorithms - so many algorithms
 - Linear Regression, Random Forest Regressor, etc.
 - Within these algorithms, many parameters

Never use testing data before anything (It will modify the model that can alter it)

- It's not ok to use test-set to compare
 - We use **cross-validation!**

N = 900
 Train = 700
 Test = 200 (Never touch it until the end)

Use different models to test using train
 and compare to the Test

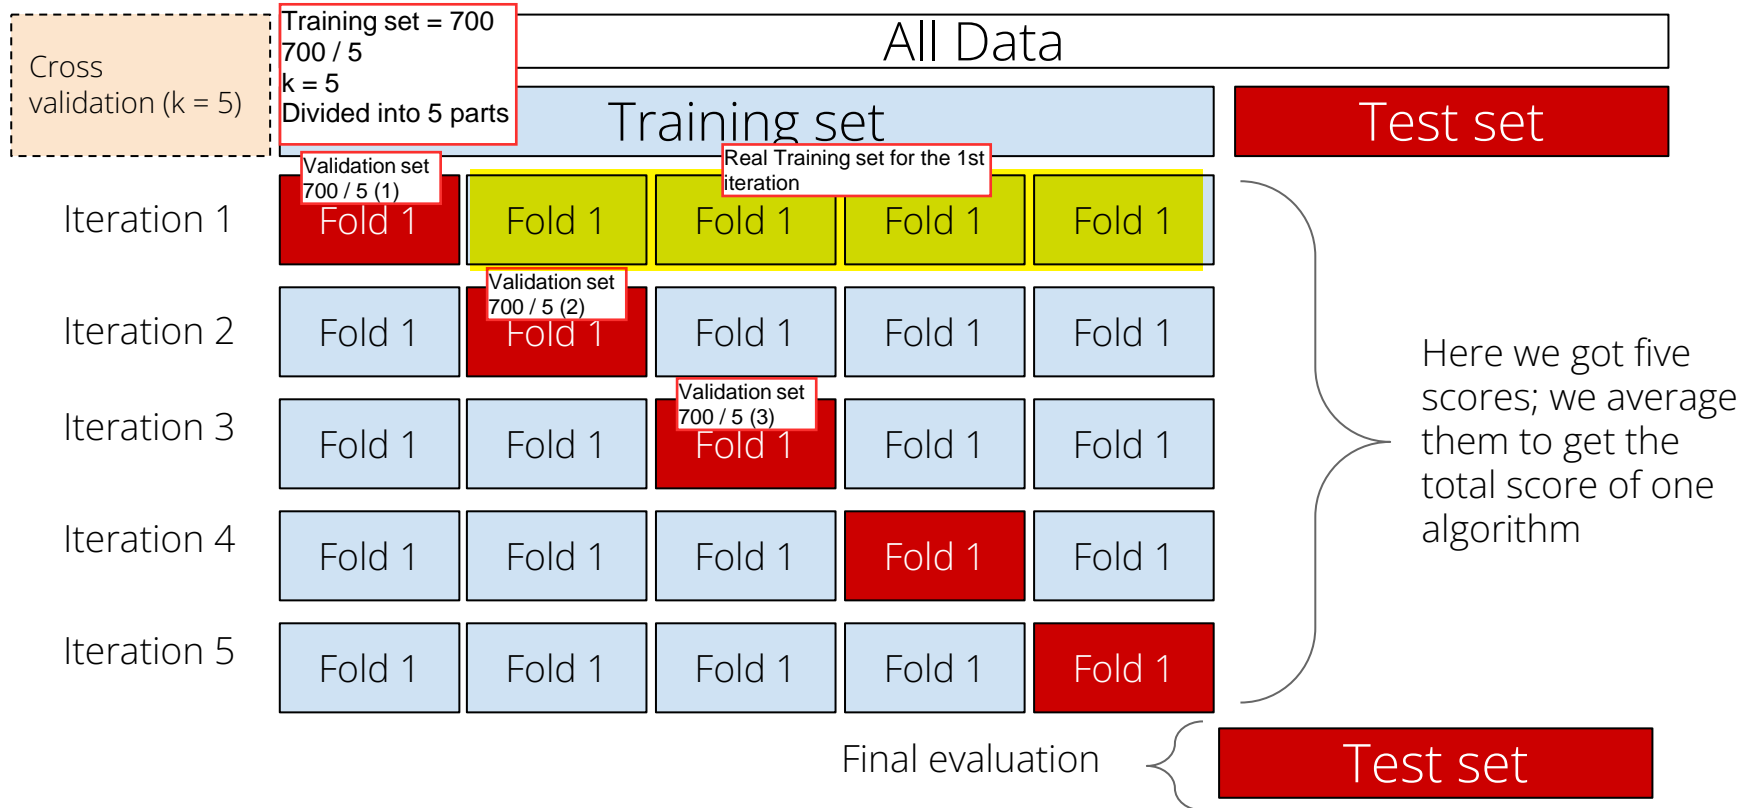
How to split?

K Fold Cross-validation

- Use no repletion method
 - Do many times & augment it
 - ite1: - ite2:
 A: 2.09 A: 3.02
 B: 1.08 B: 2.09 - Then take the mean



6. Model selection - cross validation



6. Model selection - metrics

Regression

- R^2 (R-squared) How reliable of the model compared to the mean
- MSE (Mean Squared Error)
- RMSE (Root Mean Squared Error)

$$MSE = \frac{1}{n} \sum \underbrace{\left(y - \hat{y} \right)^2}_{\substack{\text{The square of the difference} \\ \text{between actual and} \\ \text{predicted}}}$$

R-Squared range from worst to best : (-..., 1)


How it matter?

-... is worst

0 equals to the mean

1 is much better than the mean

- It doesn't mean it's good



$$R^2 = 1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}$$



6. Model selection - metrics

y	y^{\wedge}
3	1
-0.5	2
5	5
-2	-3
1	0

MSE =

Square root MSE =

R^2 =



6. Model selection - grid search

- Once we know which algorithm is best, we can further do cross-validation on that only algorithm with different parameters to find the best version
- This is called “**Grid Search**”

```
learning rate = [0.01, 0.05, 0.1, 0.5]  
algorithm = [linear reg, naive bayes]  
max_iter = [10, 50, 100, 500]
```

What combination of these hyperparameters give the most accurate results?

Solution: Use 'Grid Search'



7, 8, 9, 10: Test, Analysis, Inference, Deploy

Test - test your model with the test set

Analysis - analyze your model, e.g., feature importance, where error comes from

Inference - test your model with some unseen data and see whether it makes sense

Deploy - deploy to web using Django, Flask, FastAPI, Streamlit, etc.

